



Lessons Learned Report

Name of submission: Lessons Learned Report

Group name: Testbench for Bit-Serial RISC-V processor in SoI for Venus

Date: 2026/01/09

Organization: KTH

Authors:

Name	Email
Chinmay Purandare	chinmayp@kth.se
Lorenzo Parata	parata@ug.kth.se
Chongnan Wang	chongnan@kth.se
Xiaotian Jiang	xjia@kth.se
Xinxin Qin	xinxinq@kth.se
Paul Martin	paulmart@kth.se

Department of Electronics
KTH Royal Institute of Technology
Stockholm, Sweden

Contents

1	General summary of the project	2
2	Fulfillment of our goals	2
2.1	Follow-up of objectives	2
2.1.1	List of obtained goals	2
3	What we learned (main part)	3
3.1	Good experiences	3
3.1.1	Time Management and Scheduling	3
3.1.2	Testbench	3
3.1.3	Successful Interface Integration	4
3.1.4	Agile Project Management and Risk Awareness	4
3.1.5	Effective Collaboration Despite Physical Constraints	4
3.2	Bad Experiences	4
3.2.1	Difficulty in Project Initiation and Problem Understanding	4
3.2.2	Team Structure and Work Allocation	4
3.2.3	Dependency on Physical Hardware Availability	5
3.2.4	Limited Team Cohesion and Shared Working Time	5
3.2.5	The Cost of Assumptions	5
4	Conclusion	5

1. General summary of the project

This project focused on the implementation and verification of a parallel and a bit-serial version of a RISC-V processor. Both processor models needed to be implemented on one FPGA and tested by an advanced testbench implemented on another FPGA, with the CPU and testbench communicating through physical wires. This setup is intended to operate in the extreme environments of Venus, specifically under high-temperature conditions. The work uses KTH's experimental Silicon-on-Insulator technology, which can operate at temperatures approaching Venus's 480°C requirement. However, the test station has a constraint of only 12 pads available, requiring the development of a serialization system and communication protocol.

The project was carried out by dividing the work into two main sub-projects: the testbench development with test set and implementation, and the communication system with serialization and protocol for data transfer. This separation enabled parallel development and allowed team members to focus on different topics and architectures, ultimately achieving better efficiency. The team was organized into two subgroups: three members worked on the testbench (functional testing, CoreMark generation, and integration with RAM and UART), while three members focused on the bit-serial interface (frequency and pin properties, communication protocol specification, and RTL implementation). A significant portion of the work involved learning and adapting to unfamiliar tools, technologies, and design constraints, which strongly influenced both the schedule and working methods.

Regarding planned versus actual outcomes, the project timeline remained largely on track despite several technical challenges. The buffer time allocated during initial planning proved valuable in absorbing unforeseen delays. From a budget perspective, as an academic project no formal budget was involved, though minor unplanned expenses occurred for purchasing power adapters and communication cables between FPGA boards, which team members handled individually. In terms of results, the project successfully delivered a functional dual-FPGA system with a CPU and testbench connected by bit-serial communication. The test for CISC-V and testbench across 2 FPGAs runs successfully with measured performance at 100,198 IPS approaching the 1MHz clock of the CISC-V. The RISC-V has also been tested with the testbench, but only in simulation without the bit-serial communication.

Several significant problems impacted the project timeline and workflow. The initial codebase delivery was delayed, and upon receipt contained various compilation errors and bugs that were assumed to be outside the project scope and already corrected. This required the team to spend considerable time debugging and refactoring code that should have been functional. Additionally, some FPGA boards malfunctioned and needed replacement, and students had to purchase their own adapters and cables for inter-board communication. While these hardware issues were relatively minor and easily resolved, they introduced delays in the integration phase. The Christmas holiday period presented organizational challenges, as not all team members could work together in person. This necessitated a division of tasks between those working remotely and those with physical access to the two FPGA boards required for final integration. Since each member initially had one board, those who traveled left their boards with team members remaining on campus to enable dual-board testing to continue. The holiday period also reduced overall team availability, requiring careful task prioritization and remote workflow coordination. Despite these obstacles, the project maintained progress through adaptive planning, clear task division, and effective cross-team coordination.

2. Fulfillment of our goals

2.1. Follow-up of objectives

The project successfully achieved the majority of its stated technical and business objectives, demonstrating a functional dual-FPGA validation platform despite encountering significant implementation challenges.

2.1.1 List of obtained goals

The following table summarizes the achievement status of all project objectives defined in the initial planning phase:

ID	Objective	Status	Notes
T1	Validation of both processor architectures using a suite of at least 10 assembly-level test programs to verify core ISA functionality	Partially achieved	Both processors validated in simulation; only CISC-V validated on hardware
T2	Execution and verification of a complex benchmark algorithm (CoreMark), with results transmitted to the host PC	Achieved	Successfully executed in simulation for both processors
T3	Demonstration of stable data transfer exceeding 1,000 instructions without corruption between the two FPGAs	Achieved	Verified during CISC-V hardware testing with measured performance of 100,198 IPS
T4	(Optional) Collection of quantitative performance metrics such as maximum clock frequency, benchmark cycle counts, and FPGA resource utilization	Partially achieved	Maximum clock frequency measured; comprehensive resource analysis incomplete
B1	Complete all design, implementation, and validation phases within the planned 16-week schedule	Achieved	All major deadlines met despite technical challenges
B2	Strengthen teamwork, communication, and project management skills	Achieved	Effective collaboration and adaptive planning throughout the project
B3	Deliver a high-quality, fully functional system and comprehensive documentation	Achieved	Functional dual-FPGA system delivered with documentation

Table 1: Summary of project objectives and achievement status

The primary unfulfilled objective was the hardware validation of the RISC-V processor on the dual-FPGA platform. While the RISC-V design passed all functional tests in simulation and was verified to implement the instruction set correctly, integration with the bit-serial communication interface was not completed. This limitation stemmed directly from discrepancies in the initial codebase delivery. The project planning assumed both processor implementations would be provided as consistent black-box modules with identical external interfaces, enabling straightforward integration with the testbench through a common protocol. However, the delivered implementations featured different interface specifications, requiring significant additional integration effort that could not be accommodated within the project timeline after addressing codebase debugging needs.

The successful completion of the CISC-V validation demonstrates that the project methodology is sound and that the remaining objectives are achievable with focused effort on interface standardization and systematic performance characterization.

3. What we learned (main part)

3.1. Good experiences

3.1.1 Time Management and Scheduling

Our project timeline was not average, it was a "slow start" phase followed by high-intensity execution. The initial work was hindered by the steep learning curve of the legacy codebase—specifically the mix of VHDL and ProGram models—and logistical delays in acquiring necessary FPGA hardware. Furthermore, reliance on external feedback for complex toolchains occasionally led to stagnation, highlighting the critical need for proactive contingency planning when guidance is delayed.

Despite these early setbacks, the team successfully adapted by using the mid-term review as a catalyst to overcome inertia and accelerate development. We established a rigorous weekly meeting mechanism to synchronize the disparate timelines of processor implementation and testbench design. Additionally, individual time management played a key role; members compensated for planned absences with periods of intense workload to ensure fairness and adherence to final delivery milestones. This experience taught us that successful time management in experimental projects requires not just scheduling, but the agility to transition from stagnation to crisis management when faced with technological uncertainty.

3.1.2 Testbench

The development of the test platform provided important insights into methodologies for verifying hardware under extreme environmental conditions. A key technical lesson learned was the effectiveness of maintaining a strict "black box" level of abstraction throughout the project. By designing the test environment to interact only through the processor's

top-level instructions and data interfaces, rather than relying on internal microarchitectural signals, we ensured good portability between simulation and the final hardware implementation.

A key decision in our verification strategy was choosing a streamlined version of the CoreMark benchmark as the primary test workload. We found that simple assembly-level sanity checks are often insufficient to detect subtle data path errors. CoreMark was chosen because it provides a realistic combination of integer arithmetic, control flow, and pointer-based data structures, while being small enough for efficient simulation. This choice played a crucial role in the differential verification phase, where we compared simulation logs with the "golden logs" generated by the Spike architecture simulator. Specifically, the benchmark's use of byte and half-word memory operations exposed a critical flaw in the processor's load data path regarding sign extension—an error that simpler test programs failed to trigger.

Also, the transition from simulation to synthesizable platforms highlights the stringent limitations of physical hardware. This process underscores the importance of a design verification environment that separates behavioral transaction logic from the physical transport layer, enabling us to adapt the system from parallel bus simulation to a high-latency bit-serial hardware interface without redesigning the core logic.

3.1.3 Successful Interface Integration

One of the most rewarding experiences of the project was the moment when individual components were finally integrated together. The final integration was more complex and time-consuming than initially expected. So it is even more satisfying to see 2 interfaces of different people in the group join together. Seeing the result of our work is nice and it confirmed that the architectural decisions, communication protocols, and verification strategies were good choices.

3.1.4 Agile Project Management and Risk Awareness

One of the most positive aspects of the project was the adoption of an agile project management approach combined with structured documentation practices. Weekly sprint planning meetings and regular technical discussions with the supervisor enabled the team to continuously reassess priorities and adapt to emerging challenges.

In particular, the early definition of a risk analysis proved to be extremely valuable. Several identified risks, such as uncertainties related to legacy code quality and integration complexity, materialized during the project. Although some risks were partially underestimated in terms of effort, having them formally documented allowed the team to react quickly and reallocate time without jeopardizing major milestones. This experience demonstrated that agile planning and explicit risk awareness are essential for managing complex engineering projects under uncertainty.

3.1.5 Effective Collaboration Despite Physical Constraints

Another positive experience was the team's ability to maintain effective collaboration despite significant physical and logistical constraints. The project required access to specific FPGA hardware, and during the Christmas holiday period not all team members were able to remain on campus.

The team successfully adapted by clearly separating tasks that required physical hardware access from those that could be carried out remotely, such as simulation, documentation, and protocol development. This division, combined with frequent online communication, allowed continuous progress even when the team was geographically distributed. This experience highlighted the importance of adaptability and clear task definition when working in hybrid or partially remote engineering teams.

3.2. Bad Experiences

3.2.1 Difficulty in Project Initiation and Problem Understanding

A major challenge at the beginning of the project was the difficulty in clearly understanding the scope of the work and how to effectively start the implementation. Indeed, the project continued previous work and we had to understand what was already done and what needed to be done. Moreover, the project had existing designs and architectures that were not documented. This led to inefficiencies at the beginning. As a result, a lot of time was spent clarifying objectives, understanding code, and identifying components. This experience highlighted the importance of having a good general overview of the project before starting implementation. It can be really challenging to transition from theoretical knowledge to practice on a real project.

3.2.2 Team Structure and Work Allocation

Regarding the division of labor, we initially adopted a functional structure that split our six-person team into two equal groups: three members dedicated to the test platform, and three responsible for communication, synthesis, and clock management. While this allocation seemed logical in theory, it proved less efficient in practice. The test platform workload was not constant, leading to periods of overlap and diminishing returns for the three assigned members.

Afterward, we learned that a more flexible structure would have been better. Instead of a rigid 3-3 split, assigning two members to the test platform and make the third as a "mobile" support engineer would have been more effective. This role would have allowed for dynamic resource shifting to assist colleagues on the synthesis and clock generation tasks, who often had to handle high-pressure components in isolation. This experience underscored the importance of flexible team roles over static functional assignments to ensure a balanced workload across all project phases.

3.2.3 Dependency on Physical Hardware Availability

Another challenge was the strong dependency on physical hardware availability for final integration and validation. Since only a limited number of FPGA boards and accessories were available, delays in acquiring cables, adapters, or replacement boards directly affected progress.

This dependency reduced flexibility during critical phases of the project and increased pressure on team members who had physical access to the hardware. The experience revealed the importance of planning redundancy or alternative testing strategies when hardware resources are limited.

3.2.4 Limited Team Cohesion and Shared Working Time

One negative aspect of the project was the limited opportunity to develop strong team cohesion. Due to the nature of the project, which did not require daily physical presence, and the differing schedules of team members, much of the work was carried out individually or in small subgroups.

While this structure allowed efficient parallel progress, it reduced the amount of shared problem-solving and informal knowledge exchange within the full team. As a result, some integration issues could have been identified earlier with more frequent collective working sessions. This experience highlighted that effective teamwork is not only about task completion, but also about shared presence, continuous interaction, and collective ownership of the system.

3.2.5 The Cost of Assumptions

We assumed the delivered processor codebases would be functional and have identical interfaces. Both assumptions proved wrong. The code contained compilation errors and bugs requiring substantial debugging, while the different interface specifications prevented RISC-V hardware validation within our timeline.

Lesson: Always include acceptance testing of third-party deliverables as an early-phase activity with contingency time allocated.

4. Conclusion

This project provided a comprehensive learning experience that went far beyond the technical implementation of a processor verification platform. Working on a complex dual-FPGA system under strict physical constraints required not only solid engineering skills, but also adaptability, communication, and effective project management. The combination of legacy code, limited documentation, hardware dependency, and partial team availability made the project challenging, but also highly instructive.

From a technical perspective, the project strengthened our understanding of hardware verification methodologies, interface abstraction, and the transition from simulation to real hardware. The successful hardware validation of the CISC-V processor demonstrated the effectiveness of the adopted testbench strategy and communication protocol, while the incomplete hardware integration of the RISC-V processor highlighted the importance of early interface standardization and realistic assumptions about external code quality.

Equally important were the non-technical lessons learned. The experience emphasized that project success depends not only on individual competence, but also on flexible team structures, shared working time, and efficient communication. Issues related to task allocation, coordination overhead, and limited team cohesion revealed how easily organizational factors can become bottlenecks in technically sound projects. At the same time, adaptive planning, risk awareness, and clear prioritization proved essential in maintaining progress despite unforeseen obstacles.

Looking forward, this experience has changed how we will approach future engineering projects. We will place greater emphasis on early project scoping, interface definition, and documentation before committing to implementation. We will also advocate for more flexible team roles, increased collective working sessions during critical phases, and proactive communication strategies to reduce coordination friction. Finally, we have learned to treat risk analysis not as a formality, but as an active tool to guide decision-making throughout the project lifecycle.

In conclusion, while not all objectives were fully achieved, the project met its primary goals and delivered a functional and meaningful result. More importantly, it provided valuable insights into both technical system design and collaborative engineering practice, forming a strong foundation for future academic and professional work.